

# File Organisations and Indexes

# File organisations

- Three types of file organisations
  - Heap, Sorted and Hashed
- The COST of processing DB queries
  - SCAN, Search (Single, Range), Insert, Delete

# File Organisations

- A bit of recall
  - The basic store unit on disk (in memory) is block (page)
  - We will use page/block interchangeably.
  - One page consists of multiple data records.
- Three types
  - Heap Files (The simplest)
    - Page after page, always insert at the end
  - Sorted Files
    - Records are sorted (within and among pages) w.r.t. the search key
  - Hashed Files.
    - The pages are assigned to multiple buckets, each containing several pages
    - Each page has been assigned with an ID (bucket ID) to tell where to find the page

# COST MODEL

- Notations

- **B**: the number of pages (blocks)
- **R**: the average number of records in a page.
- **D**: average time to read or write a disk page.
- **C**: average time to process a record.
- **H**: the time required to apply a hash function to a record.

COST = Time of Execution

# Operations to be investigated

- **Scan**: fetch all records in a table.
- Search with equality selection (**SWES**)
  - “Find the students record with sid = z3087251”
- Search with Range Selection (**SWRS**)
  - “Find all students with name alphabetically after ‘Smith’”
- **Insert**: Insert a given record into the table.
- **Delete**: Delete a record with given rid.

Below, we examine the costs of these operations with respect to the 3 different file organisations.

# Heap Files

- Scan:
  - $B(D + RC)$
- SWES:
  - $0.5B(D + RC)$  on average if the selection is specified on a key.
  - Otherwise  $B(D + RC)$ .

**B**: the number of pages

**R**: the average number of records in a page (block).

**D**: average time to read or write a disk page.

**C**: average time to process a record.

# Heap Files

- SWRS:
  - $B(D + RC)$
- Insert:
  - $(D + C) + D = 2D + C$  (Always insert to the end of the file)
- Delete:
  - Only one record is involved.
    - The average cost is  $0.5B(D + RC) + D$  if rid is not given;
    - otherwise  $(D + C) + D$ .
  - Several records are involved. Expensive.

**B**: the number of pages

**R**: the average number of records in a page (block).

**D**: average time to read or write a disk page.

**C**: average time to process a record.

# Sorted Files

Sorted on a search key

(a combination of one or more fields)

Querying against the search key:

- Scan:  $B(D + RC)$ .
- SWES:
  - $O(D \log_2 B + C \log_2 R)$  if single record. (log comes from binary search)
  - $O(D \log_2 B + C \log_2 R + \text{\#matches})$ .
- SWRS:  $O(D \log_2 B + C \log_2 R + \text{\#matches})$ .
- Insert: expensive.
  - Search cost plus  $2 * (0.5B(D + RC))$ .
- Delete: expensive.
  - Search cost plus  $2 * (0.5B(D + RC))$ .



# Hashed Files

- The pages in a file are grouped into buckets.
- The buckets are defined by a hash function.
- Pages are kept at about 80% occupancy.

Assume the data manipulation is based on the hash key.

- Scan:  $1.25B(D + RC)$ .
- SWES:  $H + D + 0.5RC$  if each hash bucket contains only one page.
- SWRS:  $1.25B(D + RC)$ . (No help from the hash structure)
- Insert: Search cost plus  $C + D$  if one block involved.
- Delete: Search cost plus  $C + D$  if one block involved.

# Summary

File Type	Scan	Equality Search	Range Search	Insert	Delete
Heap	BD	0.5 BD	BD	Search + D	Search + D
Sorted	BD	D logB	$D \log B + \text{\# matches}$	Search + BD	Search + BD
Hashed	1.25 BD	D	1.25 BD	2 D	Search + BD

## A Comparison of I/O Costs

(I/O Costs: The cost of reading from or writing to disk)

# Introduction to Index

- What is index?
- What is the basic organisation of index?

# What is index?



Book has been arranged via categories, subjects.

Same categories will be stored in the same area.

Book is data, then  
Catalogue is the index

# What is index?



Phone number has been arranged in Alphabets, groups (work, friends, ...)

Phone number is the data  
The Alphabets and groups are the index.

# What is index?

- Auxiliary data
- Properly organised (data structure)
- To facilitate data search

# Indexes


Word indexes in a book:

## INDEXES

aardvark 25,36	lion . . . . . 18
bat . . . . . 12	llama 17,21,22
cat . . . . 1,5,12	sloth . . . . . 18
dog . . . . . 3	tiger . . . . . 18
elephant . . 17	wombat . . . 27
emu . . . . . 28	zebra . . . . . 19

- A table of key values, where each entry gives places where key is used.
- Aim: efficient access to records via key.

# Indexing Structure






# Indexing Structure

- Index is collection of data entries  $k^*$ :
  - Each data entry  $k^*$  contains enough information to retrieve (one or more) records with search key value  $k$ .
- Indexing:
  - How are data entries organised in order to support efficient retrieval of data entries with a given search key value?
  - What is stored as a data entry?


# Clustered Index

- Clustered: Sort and store the data records in the table or view based on the search keys of the index.
- Typically, the search key of file is the same as the search key of index.

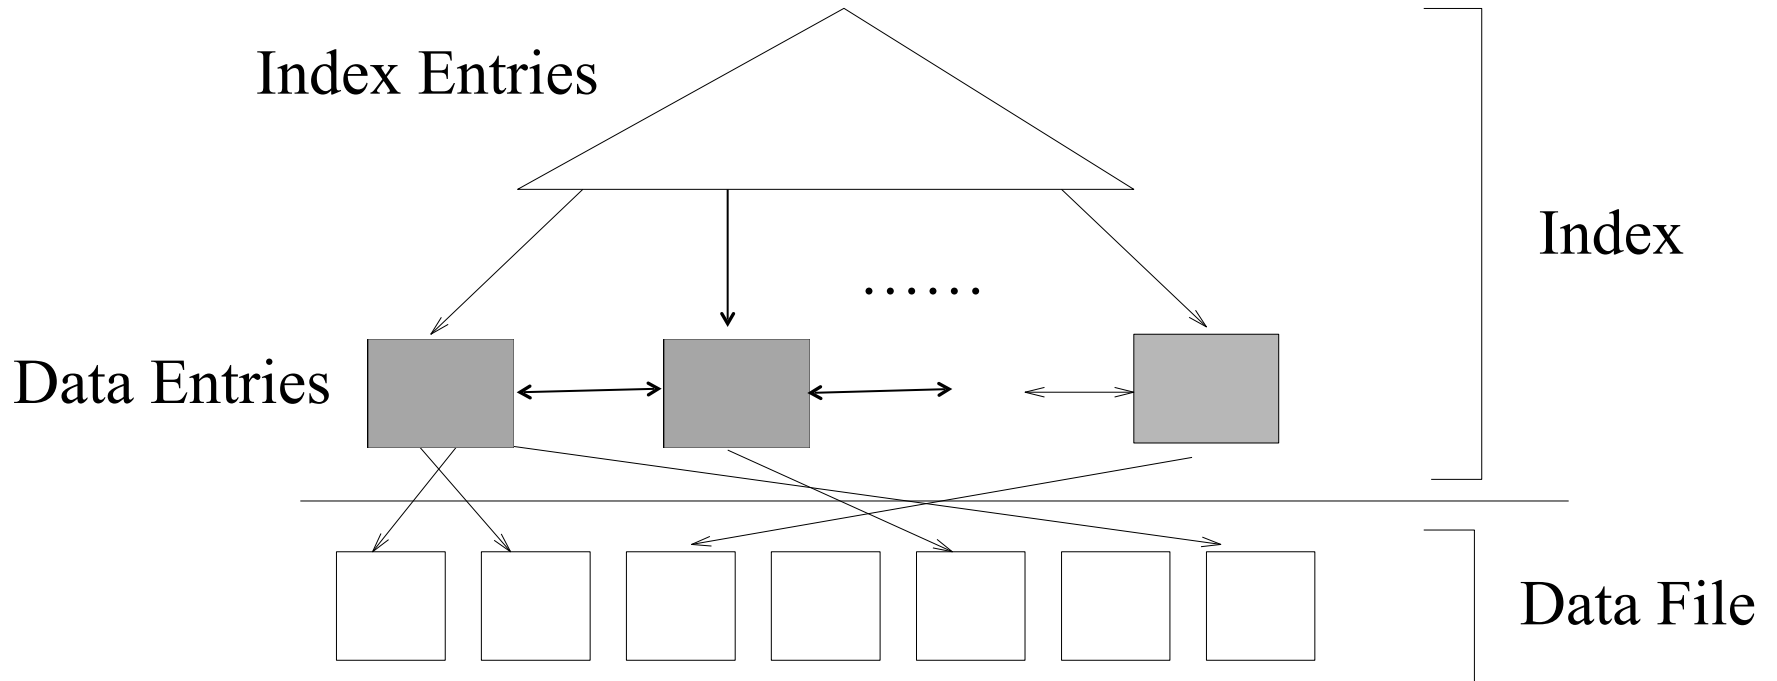


# Clustered Index

- Clustered indexes are relatively expensive to maintain.
- A data file can be clustered (indexed) on **at most one** search key.



# Unclustered Index

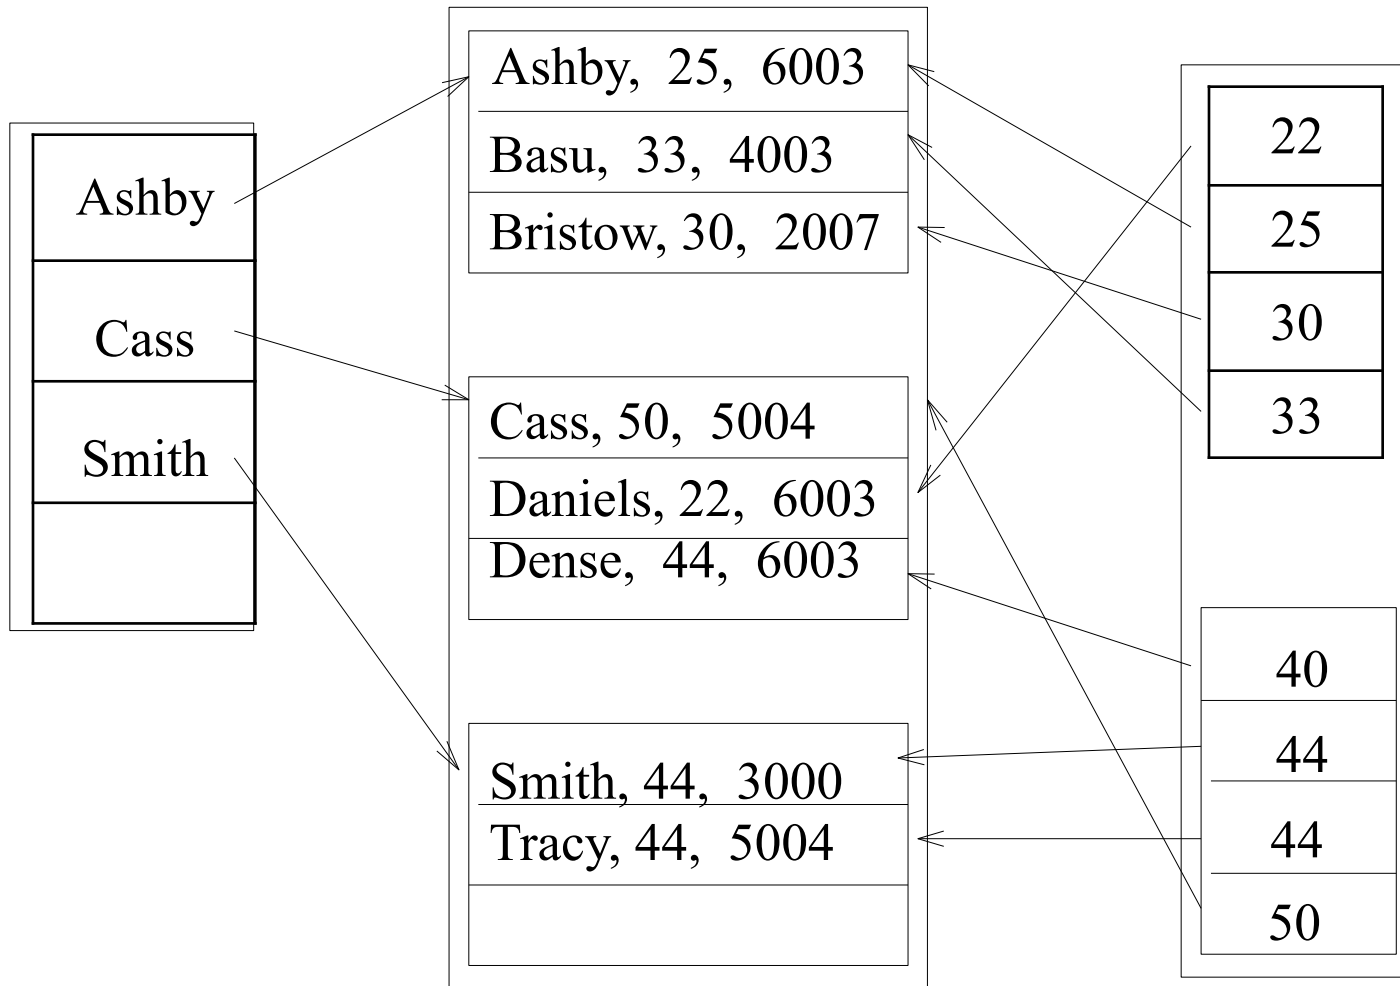


# Dense VS Sparse Indexes

- Dense Index and Sparse Index
  - Dense Index contains (**at least**) one data entry for every search key value.
  - Sparse Index: one search key can point to a set of data entries

Q: Can we build a sparse index that is not clustered?

# Dense VS Sparse Indexes



Sparse Index

Dense Index

# Primary and Secondary Indexes

- Primary: Indexing fields include **primary key**.
  - There can be at most one primary index for a table
- Secondary: otherwise.
  - Composite search keys: search key contains several fields.