

Transaction Management

Air-line Reservation

- 10 available seats vs 15 travel agents.
- How do you design a robust and fair reservation system?
 - Do not enough resources
 - Fair policy to every body
 - Robustness

Failures

Number of factors might cause failures in user requirements processing.

1. System failure:
 - Disk failure - e.g. head crash, media fault.
 - System crash - unexpected failure requiring a reboot.
2. Program error - e.g. a divide by zero.
3. Exception conditions - e.g. no seats for your reservation.
4. Concurrency control - e.g. deadlock, expired locks.

To handle failures correctly and efficiently

Each database user must express his requirements as a set of program units.

Each program unit is a *transaction* that either

- accesses the contents of the database, or
- changes the state of the database, from one consistent state to another.

- Sydney → Tokyo → LA
→ N.Y
- It does not make sense
only partial trip has tickets

Example transaction: buy a ticket from Sydney to N.Y. by JAL.

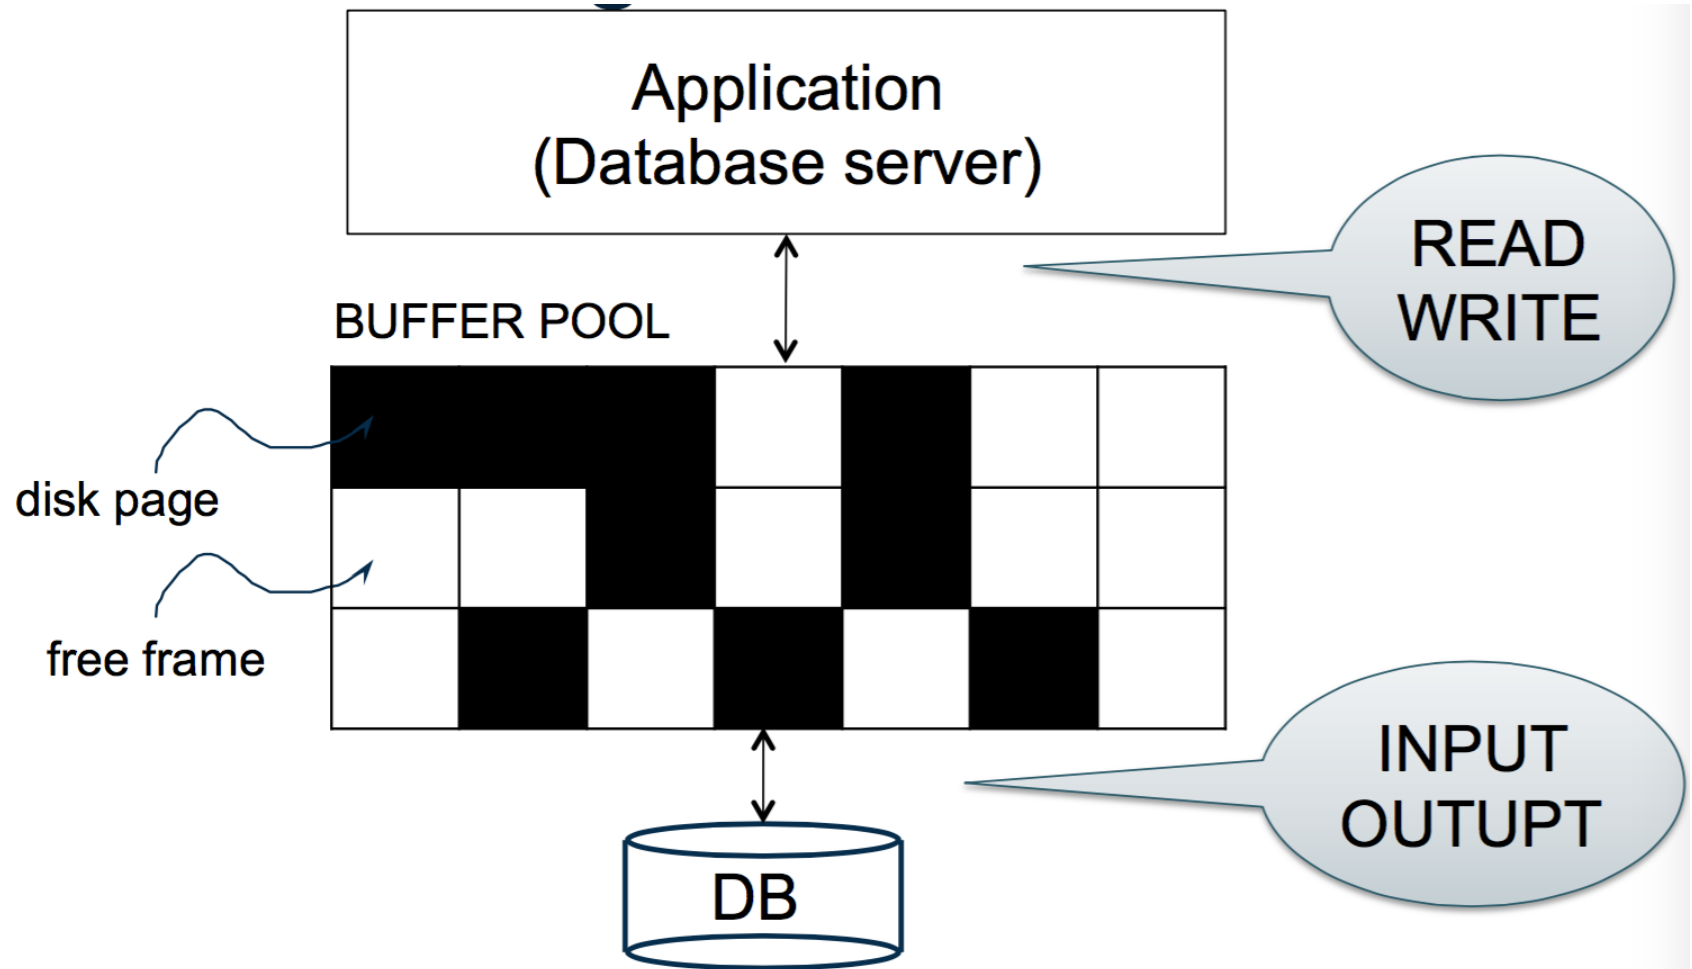
A transaction must be treated as an *atomic* unit.

Transaction Processing

Three kinds of operations may be used in a transaction:

- *Read.*
- *Write.*
- Computation.

Buffer Management in a DBMS



Read

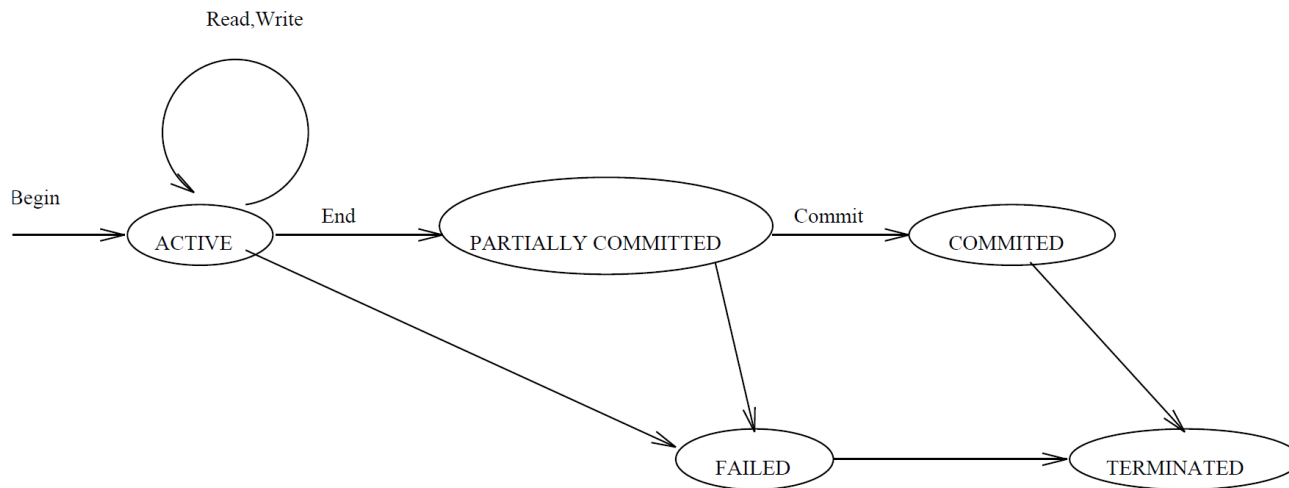
1. Compute the data block that contains the item to be read
2. Either
 - find a buffer containing the block, or
 - read from disk into a buffer
3. Copy the value from the buffer.

Write

1. Compute the disk block containing the item to be written,
2. Either
 - find a buffer containing the block, or
 - read from disk into a buffer,
3. Copy the new value into the buffer,
4. At some point (maybe later), write the buffer back to disk.

Processing States of a Transaction

- The typical processing states are illustrated in the figure below (E/N Fig 17.4):



- Partially committed point:*** At this point, check and enforce the correctness of the concurrent execution.
- Committed state:*** Once a transaction enters the committed state, it has concluded its execution successfully.

Desirable Properties of Transaction Processing **ACID**

- Atomicity: A transaction is either performed in its entirety or not performed at all.
- Consistency preservation: A correct execution of the transaction must take the database from one consistent state to another.
- Isolation: A transaction should not make its updates visible to other transactions until it is committed.
- Durability or permanency: Once a transaction changes the database and the changes are committed, these changes must never be lost because of subsequent failure.

Problems without Enforcing ACID

- For a banking system,
 - If durability is not enforced, then a customer may lose a deposit.
 - If consistency preservation is not enforced, then the bank runs a high risk of bankrupt. E.g., run-over upper-limit.
- Below are the problems if atomicity and isolation are not enforced in a concurrent execution of transactions.

Lost Update Problem (Isolation is not enforced)

- Suppose we have these two transactions, T_1 and T_2 :

T_1 :

```
read(X)
 $X \leftarrow X + N$ 
write(X)
read(Y)
 $Y \leftarrow Y - N$ 
write(Y)
```

T_2 :

```
read(X)
 $X \leftarrow X + M$ 
write(X)
```

- Let us see what may happen if T_1 and T_2 are executed concurrently in an uncontrolled way:

Suppose initially that $X = 100$; $Y = 50$; $N = 5$ and $M = 8$.

Database	T_1	T_2
$X = 100, Y = 50$	$X = ?, Y = ?$	$X = ?$
$X = 100, Y = 50$	read(X) $X = 100, Y = ?$	$X = ?$
$X = 100, Y = 50$	$X \leftarrow X + N$ $X = 105, Y = ?$	$X = ?$
$X = 100, Y = 50$	$X = 105, Y = ?$	read(X) $X = 100$
$X = 100, Y = 50$	$X = 105, Y = ?$	$X \leftarrow X + M$ $X = 108$
$X = 105, Y = 50$	write(X) $X = 105, Y = ?$	$X = 108$
$X = 105, Y = 50$	read(Y) $X = 105, Y = 50$	$X = 108$
$X = 108, Y = 50$	$X = 105, Y = 50$	write(X) $X = 108$
$X = 108, Y = 50$	$Y \leftarrow Y - N$ $X = 105, Y = 45$	$X = 108$
$X = 108, Y = 45$	write(Y) $X = 105, Y = 45$	$X = 108$

- At the end of T_1 and T_2 , X should be 113, Y should be 45.
- The update $X \leftarrow X + N$ has been lost.

The Temporary Update Problem

Database	T_1	T_2
$X = 100, Y = 50$	$X = ?, Y = ?$	$X = ?$
$X = 100, Y = 50$	read(X) $X = 100, Y = ?$	$X = ?$
$X = 100, Y = 50$	$X \leftarrow X + N$ $X = 105, Y = ?$	$X = ?$
$X = 105, Y = 50$	write(X) $X = 105, Y = ?$	$X = ?$
	FAILS	
$X = 105, Y = 50$		read(X) $X = 105$
$X = 105, Y = 50$		$X \leftarrow X + M$ $X = 113$

Recover from the disk



Several possibilities for what might happen next:

Database	T_1	T_2
X = 105, Y = 50		X = 113
X= 100, Y = 50	Case 1: DBMS undoes T_1	
X=113, Y=50		Write (X) X= 113

Database	T_1	T_2
X = 105, Y = 50		X = 113
X= 105, Y = 50	Case 2: DBMS does nothing to T_1	
X=113, Y=50		Write (X) X= 113

Database	T_1	T_2
X = 105, Y = 50		X = 113
X= 105, Y = 50		X = 113
X=100, Y=50	Case 3: DBMS undoes T_1	
		Write (X), X= 113 X = 100

Case 1:

Database	T_1	T_2
$X = 105, Y = 50$		$X = 113$
DBMS undoes T_1		
$X = 100, Y = 50$		$X = 113$
$X = 113, Y = 50$		write(X) $X = 113$

Case 2:

Database	T_1	T_2
$X = 105, Y = 50$		$X = 113$
DBMS does nothing about T_1		
$X = 105, Y = 50$		$X = 113$
$X = 113, Y = 50$		write(X) $X = 113$

Case 3:

Database	T_1	T_2
$X = 105, Y = 50$		$X = 113$
$X = 113, Y = 50$		write(X) $X = 113$
DBMS undoes T_1		
$X = 100, Y = 50$		$X = 113$

- In case 1 and 2, only half of T_1 has been executed.
- In case 3, T_2 has been lost.

The Incorrect Summary Problem

T_1	T_3
	$sum \leftarrow 0$
	read(A)
	$sum \leftarrow sum + A$
	\vdots
read(X)	
$X \leftarrow X - N$	
write(X)	
	\vdots
	read(X)
	$sum \leftarrow sum + X$
	read(Y)
	$sum \leftarrow sum + Y$
	\vdots
read(Y)	
$Y \leftarrow Y + N$	
write(Y)	
	\vdots

- Here the sum calculated by T_3 will be wrong by N .

Recover from Failures

- Ensure the **A** in **ACID**
- Log-based Recovery
 - Undo logging
 - Redo logging
 - Undo/Redo logging

System Log

- System Log
 - The system needs to record the states information to recover failures correctly.
 - The information is maintained in a log (also called journal or audit trail).
 - The system log is kept in hard disk but maintains its current contents in main memory.

System Log

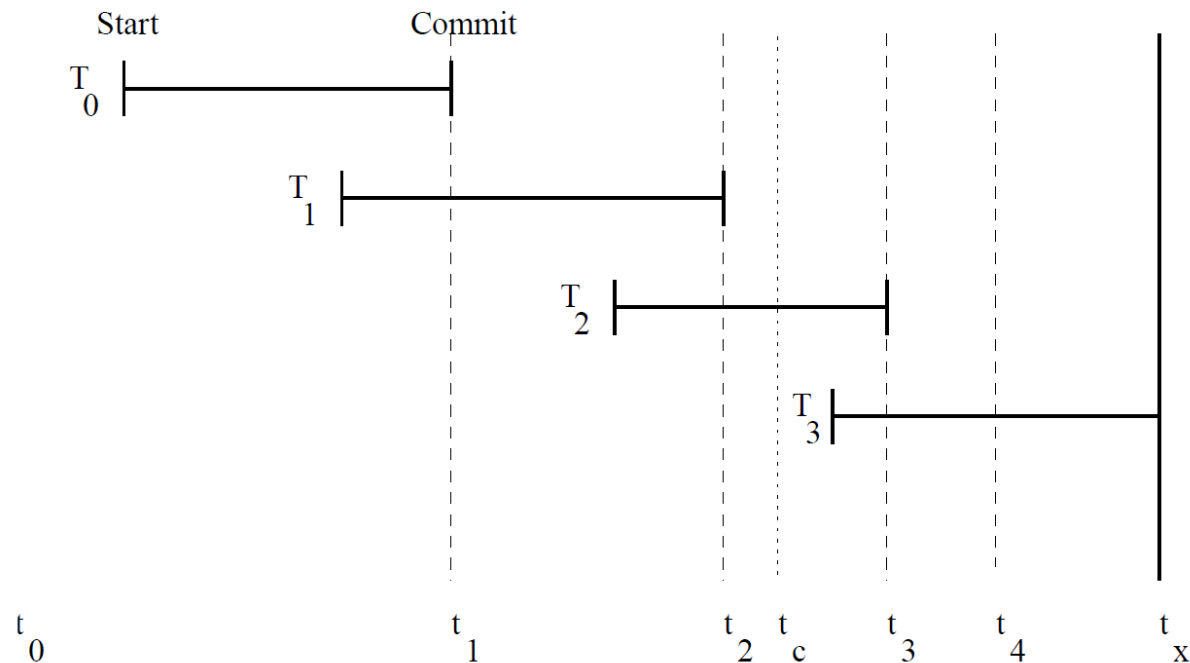
- Start transaction marker [start transaction, T]: Records that transaction T has started execution.
- [read item, T , X]: Records that transaction T has read the value of database item X .
- [write item, T , X , old value, new value]: Records that T has changed the value of database item X from old value to new value.
- Commit transaction marker [commit, T]: Records that transaction T has completed successfully, and asserts that its effect can be committed (recorded permanently) to the database.
- [abort, T]: Records that transaction T has been aborted.

System Log (Cont'd)


- In fact some other entries (rollback, undo, redo) are also required for a recovery method.
- These entries allow the recovery manager to *rollback* an unsuccessful transaction (undo any partial updates).

Recovery

- Let us see how the log might be used to recover from a system crash.
- The diagram below shows transactions between the last system backup and a crash.




Recovery (Cont'd)




- The database on disk will be in a state somewhere between that at t_0 and the state at t_x .
- The same is also true for log entries.

Recovery (Cont'd)

- We will assume that the *write-ahead log strategy* is used. This means that
 - old data values must be force-written to the log (i.e. the buffer must be copied to disk) before any change can be made to the database, and
 - the transaction is regarded as committed when the new data values and the commit marker have been force-written to the log.
- Thus the log is force-written at least at t_1 , t_2 and t_3 in the above.




- Suppose the log was last written to disk at t_4 .
- By examining the log:
 1. We know that T_0, T_1 and T_2 have committed and their effects should be reflected in the database after recovery.
 2. But we do not know whether the effects of T_0, T_1 and T_2 were reflected at the time of the crash.
 3. We also know that T_3 has started, may have modified some data, but is not committed. Thus T_3 should be undone.



- The database can be recovered by rolling back T_3 using the old data values from the log, and redoing the changes made by $T_0 \dots T_2$ using the new data values (for these committed transactions) from the log.
- Notice that instead of rolling back, the database could have been restored from the backup. This might be necessary in the event of a disk crash for example (for this reason, the log should be stored on an independent disk pack).

Checkpoints

- Notice also that using this system, the longer the time between crashes, the longer recovery may take.
- To avoid this problem, the system may take *checkpoints* at regular intervals.
- To do this:
 - a *start of checkpoint* marker is written to the log, then
 - the database updates in buffers are force-written, then
 - an *end of checkpoint* marker is written to the log.



- In our example, suppose a checkpoint is taken at time t_c . Then on recovery we only need redo T_2 .

Recall: Desirable Properties of Transaction Processing: **ACID**

- Atomicity: A transaction is either performed in its entirety or not performed at all.
- Consistency preservation: A correct execution of the transaction must take the database from one consistent state to another.
- Isolation: A transaction should not make its updates visible to other transactions until it is committed.
- Durability or permanency: Once a transaction changes the database and the changes are committed, these changes must never be lost because of subsequent failure.

Concurrency Control

- Multiple concurrent transactions T_1, T_2, \dots
- They read/write common elements A_1, A_2, \dots
- How can we prevent unwanted interference?

The Scheduler is responsible for that

Schedules of Transactions

- To fully utilise resources, desirable to interleave the operations of transactions in an appropriate way.
- For example, if one transaction is waiting for I/O to complete, another transaction can use the CPU.
- A *schedule* S of the transactions T_1, \dots, T_n
 - is a sequential ordering of the operations of T_1, \dots, T_n , and
 - preserves the ordering of operations in each transaction T_i .

Example Schedules

(a)		(b)	
T_1	T_2	T_1	T_2
read(X) $X \leftarrow X + N$ write(X) read(Y) $Y \leftarrow Y - N$ write(Y)	 read(X) $X \leftarrow X + M$ write(X)	 read(X) $X \leftarrow X + N$ write(X) read(Y) $Y \leftarrow Y - N$ write(Y)	read(X) $X \leftarrow X + M$ write(X)

Example Schedules (Cont.)

(c)		(d)	
T_1	T_2	T_1	T_2
read(X) $X \leftarrow X + N$		read(X) $X \leftarrow X + N$	
	read(X) $X \leftarrow X + M$	write(X)	
write(X) read(Y)			read(X) $X \leftarrow X + M$
	write(X)	read(Y) $Y \leftarrow Y - N$	write(X)
$Y \leftarrow Y - N$ write(Y)		write(Y)	

Serial Schedule

- As we have seen, if operations are interleaved arbitrarily, incorrect results may occur.
- However, it is reasonable to assume that schedules (a) and (b) in the figure will give correct results (as long as the transactions are independent).
- (a) and (b) are called *serial* schedules, and we will assume that *any serial schedule is correct*.

(a)		(b)	
T_1	T_2	T_1	T_2
read(X) $X \leftarrow X + N$ write(X) read(Y) $Y \leftarrow Y - N$ write(Y)	read(X) $X \leftarrow X + M$ write(X)	read(X) $X \leftarrow X + N$ write(X) read(Y) $Y \leftarrow Y - N$ write(Y)	read(X) $X \leftarrow X + M$ write(X)

Serializable Schedule

- Notice that schedule (d) always produces the same result as schedules (a) and (b), so it should also give correct results.
- A schedule is *serializable* if it always produces the same result as some serial schedule.
- Notice that schedule (c) is not serializable.

(c)		(d)	
T_1	T_2	T_1	T_2
read(X) $X \leftarrow X + N$		read(X) $X \leftarrow X + N$ write(X)	
write(X) read(Y)	read(X) $X \leftarrow X + M$		read(X) $X \leftarrow X + M$ write(X)
$Y \leftarrow Y - N$ write(Y)	write(X)	read(Y) $Y \leftarrow Y - N$ write(Y)	

Scheduling Transactions

- Serial schedule: Schedule that does not interleave the actions of different transactions.
- Equivalent schedules: For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- Serializable schedule: A schedule over a set S of transactions is equivalent to some serial execution of the set of committed transactions in S .

Serializability

Note: If each transaction preserves consistency, every serializable schedule preserves consistency.


Conflict Serializable Schedules

- Two schedules are *conflict equivalent* if:
 - Involve the same actions of the same transactions
 - Every pair of conflicting actions is ordered the same way
- Schedule S is *conflict serializable* if S is conflict equivalent to some serial schedule

Conflict Serializability

- Any conflict serializable schedule is also a serializable schedule (why ?)
- The inverse is not true.

Equivalent, but not
conflict equivalent



Testing Conflict Serializable

- Why not run only serial schedules? That is, run one transaction after the other?

Because of very poor throughput due to disk latency

- When there are only two transactions, there are only two serial schedules - for n transactions there will be $n!$.
- Fortunately there is an efficient algorithm to check whether a schedule is conflict serializable without checking all these possibilities.

Check Conflict Serializability


- *Algorithm*

Step 1: Construct a *schedule* (or *precedence*) graph – a *directed graph*.

Step 2: Check if the graph is *cyclic*:

- Cyclic: non-serializable.
- Acyclic: serializable.

- A *directed graph* $G = (V, A)$ consists of
 - a vertex set V , and
 - an arc set A such that each arc connects two vertices.
- G is *cyclic* if G contains a directed cycle.



Cyclic Graph

Construct a Schedule Graph $G_S = (V, A)$ for a schedule S


1. A vertex in V represents a transaction.
2. For two vertices T_i and T_j , an arc $T_i \rightarrow T_j$ is added to A if
 - there are two *conflicting* operations $O_1 \in T_i$ and $O_2 \in T_j$,
 - in S , O_1 is before O_2 .

Two operations O_1 and O_2 are *conflicting* if

- they are in different transactions but on the same data item,
- one of them must be a write.

Example 1


$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$



This schedule is conflict-serializable

Example 2


$r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B)$



This schedule is NOT conflict-serializable


Example 1:

Schedule	T_1	T_2	T_3
read(A)	read(A)		
read(B)		read(B)	
$A \leftarrow f_1(A)$	$A \leftarrow f_1(A)$		
read(C)			read(C)
$B \leftarrow f_2(B)$		$B \leftarrow f_2(B)$	
write(B)		write(B)	
$C \leftarrow f_3(C)$			$C \leftarrow f_3(C)$
write(C)			write(C)
write(A)	write(A)		
read(B)			read(B)
read(A)		read(A)	
$A \leftarrow f_4(A)$		$A \leftarrow f_4(A)$	
read(C)	read(C)		
write(A)		write(A)	
$C \leftarrow f_5(C)$	$C \leftarrow f_5(C)$		
write(C)	write(C)		
$B \leftarrow f_6(B)$			$B \leftarrow f_6(B)$
write(B)			write(B)



Example 2:

Schedule	T_1	T_2	T_3
read(A)	read(A)		
$A \leftarrow f_1(A)$	$A \leftarrow f_1(A)$		
read(C)	read(C)		
write(A)	write(A)		
$A \leftarrow f_2(C)$	$A \leftarrow f_2(C)$		
read(B)		read(B)	
write(C)	write(C)		
read(A)		read(A)	
read(C)			read(C)
$B \leftarrow f_3(B)$		$B \leftarrow f_3(B)$	
write(B)		write(B)	
$C \leftarrow f_4(C)$			$C \leftarrow f_4(C)$
read(B)			read(B)
write(C)			write(C)
$A \leftarrow f_5(A)$		$A \leftarrow f_5(A)$	
write(A)		write(A)	
$B \leftarrow f_6(B)$			$B \leftarrow f_6(B)$
write(B)			write(B)



- Unfortunately, testing for serializability on the fly is not practical.
- Instead, a number of protocols have been developed which ensure that if every transaction obeys the rules, then *every* schedule will be serializable, and thus correct.

Concurrency Control Methods

- **Locking Mechanism**

The idea of locking some data item X is to:

- give a transaction exclusive use of the data item X ,
- do not restrict the access of other data items.

This prevents one transaction from changing a data item currently being used in another transaction.

- We will discuss a simple locking scheme which locks individual items, using read and write locks

Locking Rules

- In this schema, every transaction T must obey the following rules.
- 1) If T has only one operation (read/write) manipulating an item X :
 - obtain a read lock on X before reading it,
 - obtain a write lock on X before writing it,
 - unlock X when done with it.
- 2) If T has several operations manipulating X :
 - obtain one proper lock only on X :
a read lock if all operations on X are reads;
a write lock if one of these operations on X is a write.
 - unlock X after the last operation on X in T has been executed.

Locking Rules (cont.)

- In this scheme,
 - Several read locks can be issued on the same data item at the same time.
 - A read lock and a write lock cannot be issued on the same data item at the same time, neither two write locks.
- This still does not guarantee serializability.

Example: Based on E/N Fig 18.3.

T_1	T_2
read_lock(Y) read(Y) unlock(Y)	read_lock(X) read(X) unlock(X) write_lock(Y) read(Y) $Y \leftarrow X + Y$ write(Y) unlock(Y)
write_lock(X) read(X) $X \leftarrow X + Y$ write(X) unlock(X)	

Two Phase Locking (2PL)

- To guarantee serializability, transactions must also obey the *two-phase locking protocol*:
 - Growing Phase: all locks for a transaction must be obtained before any locks are released, and
 - Shrinking Phase: gradually release all locks (once a lock is released no new locks may be requested).

Two Phase Locking (2PL) (Cont.)

Example: Based on E/N Fig 18.4.

T_1

read_lock(Y)
read(Y)
write_lock(X)
unlock(Y)
read(X)
$X \leftarrow X + Y$
write(X)
unlock(X)

- Locking thus provides a solution to the problem of correctness of schedules.

Two phase locking ensures conflict serializability

Deadlock


- A problem that arises with locking is **deadlock**.
- Deadlock occurs when two transactions are each waiting for a lock on an item held by the other.

T_1	T_2
write_lock(X) read(X)	
write_lock(Y) **waiting for Y*** **waiting for Y***	write_lock(Y) read(Y)
	write_lock(X) ***waiting for X***

Deadlock Check

- Create the *wait-for graph* for currently active transactions:
 - create a vertex for each transaction; and
 - an arc from T_i to T_j if T_i is waiting for an item locked by T_j .
- If the graph has a cycle, then a *deadlock* has occurred.

Example:



Several methods to deal with deadlocks

- deadlock detection
 - periodically check for deadlocks, abort and rollback some transactions (restart them later). This is a good choice if transactions are very short or very independent.

Several methods to deal with deadlocks (Cont.)

- deadlock prevention - Assign priorities based on timestamps.
Assume T_i wants a lock that T_j holds. Two policies are possible:
 - Wait-Die: If T_i has higher priority, T_i waits for T_j ; otherwise T_i aborts
 - Wound-wait: If T_i has higher priority, T_j aborts; otherwise T_i waits
- If a transaction re-starts, make sure it has its original timestamp

Timestamp ordering

- The idea here is:
 - to assign each transaction a timestamp (e.g. start time of transaction), and
 - to ensure that the schedule used is equivalent to executing the transactions in timestamp order

- Each data item, X , is assigned
 - a read timestamp, $read\ TS(X)$ - the latest timestamp of a transaction that read X , and
 - a write timestamp, $write\ TS(X)$ - the latest timestamp of a transaction that write X .

- These are used in read and write operations as follows. Suppose the transaction timestamp is T .

read(X):

```
If T >= write_TS(X) then
    { execute read(X);
      if T >= read_TS(X) then
        read_TS(X) <- T }
else
    rollback the transaction and restart
```

write(X):

```
If T >= read_TS(X) and T >= write_TS(X) then
    { execute write(X); write_TS(X) <- T }
else
    rollback and restart
```

- **Thomas' write rule:**

write(X):

```
If T < read_TS(X) then
    rollback and restart
else if T < write_TS(X) then
    ignore the write
else
    { execute write(X);
      write_TS(X) <- T }
```


- Some problems:
 - Cyclic restart: There is no deadlock, but a kind of livelock can occur - some transactions may be constantly aborted and restarted.
 - Cascading rollback: When a transaction is rolled back, so are any transactions which read a value written by it, and any transactions which read a value written by them . . . etc. This can be avoided by not allowing transactions to read values written by uncommitted transactions (make them wait).

Multiversioning

- Similar to the timestamp ordering approach; but is allowed to access “old” versions of a table.
- A history of the values and timestamps (versions) of each item is kept.
- When the value of an item is needed, the system chooses a **proper** version of the item that maintains serializability.
- This results in fewer aborted transactions at the cost of greater complexity to maintain more versions of each item.

- We will look at a scheme, several versions X_1, \dots, X_k of each data item are kept. For each X_i we also keep
 - *read* $TS(X_i)$ - as for timestamp ordering.
 - *write* $TS(X_i)$ - as for timestamp ordering.

- Read and write are done as follows for a transaction P with timestamp T .

read(X):

```
Find  $X_i$  s.t.  $\text{write\_TS}(X_i)$  is the  
    highest write timestamp but  $\leq T$   
update  $\text{read\_TS}(X_i)$  (and do  $\text{read}(X_i)$ )  
return  $X_i$  as the value for  $X$ 
```

write(X):

```
Find Xi s.t. write_TS(Xi) is the
    highest write timestamp but <= T
if T < read_TS(Xi) then
    rollback and restart
else
    { create a new version X(k+1) of X;
      set read_TS(X(k+1)) to T;
      set write_TS(X(k+1)) to T}
```

- *Note:* Cascading rollback and cyclic restart problems can still occur, but should be reduced.
- However, there is an increased overhead in maintaining multiple versions of items.

Optimistic scheduling

- In two-phase locking, timestamp ordering, and multiversioning concurrency control techniques, a certain degree of checking is done **before** a database operation can be executed.
- The idea here is to push on and hope for the best!
- No checking is done while the transaction is executing.

- The protocol has three phases.
 - read phase - A transaction can read data items from the database into local variables. However, updates are applied only to local copies of the data items kept in the transaction workspace.
 - validation phase - checks are made to ensure that serializability is not violated,
 - write phase -if validation succeeds, updates are applied and the transaction is committed. Otherwise, the updates are discarded and the transaction is restarted.

- A scheme uses timestamps and keeps each transaction's
 - read-set - the set of items read by the transaction,
 - write-set - the set of items written by the transaction.
- During validation, we check that the transaction does not interfere with any transaction that is committed or currently validating.

- Each transaction T is assigned 3 timestamps:
 $Start(T)$, $Validation(T)$, $Finish(T)$.
- To pass the validation test for T , one of the following must be true:
 - 1. $Finish(S) < Start(T)$; or
 - 2. for S s.t. $Start(T) < Finish(S)$, then
 - a) write set of S is disjoint from the read set of T , and
 - b) $Finish(S) < Validation(T)$.

Optimistic control is a good option if there is not much interaction between transactions.

2PL vs. TSO vs. MV vs. OP

- A Comparison among two-phase locking (2PL), timestamp ordering (TSO), multiversioning (MV), optimistic (OP) concurrency control techniques.
- MV should provide the greatest concurrency degree (in average). However, we need to maintain multiversions for each data item.
- 2PL can offer the second greatest concurrency degree (in average); but will result in deadlocks. To resolve the deadlocks, either
 - need additional computation to detect deadlocks and to resolve the deadlocks, or
 - reduce the concurrency degree to prevent deadlocks by adding other restrictions.

2PL vs. TSO vs. MV vs. OP (cont.)

- If most transactions are very short, we can use 2PL + deadlock detection and resolution.
- TSO has a less concurrency degree than that of 2PL if a proper deadlock resolution is found. However, TSO does not cause deadlocks. Other problems, such as cyclic restart and cascading rollback, will appear in TSO.
- If there are not much interaction between transactions, OP is a very good choice. Otherwise, OP is a bad choice.